

Deep Learning for Natural Language Processing

Continuous bag-of-words representations




UNIVERSITY OF
GOTHENBURG

CHALMERS


WASP | WALLENBERG AI
AUTONOMOUS SYSTEMS
AND SOFTWARE PROGRAM

Richard Johansson

`richard.johansson@gu.se`



recap: document classification with BoW representation




the continuous bag-of-words representation

- ▶ we return to the question of **representation**
- ▶ the **continuous bag-of-words** (CBoW) or **neural bag-of-words** representation computes a sum or mean of word embeddings over a document

$$\text{CBoW}(D) = \frac{1}{L} \sum_{i=1}^L \text{embedding}(w_i)$$

building a document categorizer with CBoW



- ▶ we can combine the CBoW representation with a linear or nonlinear (neural) classifier



- ▶ when we train the complete system, the word embeddings are **learned automatically** for this task
- ▶ this is our first example of a **learned representation**: a key selling point for deep learning in NLP

famous examples of CBoW architecture

- ▶ DAN (Iyyer et al., 2015)
- ▶ fastText (Joulin et al., 2017)



discussion


what do the word embeddings represent?

what do the word embeddings represent?

- ▶ intuitively, a word embedding represents the “meaning content” of this word **with respect to the target task**

limitations of the CBoW approach

- ▶ still more or less a word-spotting approach
- ▶ we can't handle **interactions** between words: negations, . . .
- ▶ in some applications, we need a more generalized notion of “patterns”




addressing the limitations of CBoW

- ▶ next block: **pre-training** embeddings
- ▶ and then: **convolutional** and **recurrent** models

case study: BoW and CBoW classification

- ▶ in the notebook for this lecture, we will implement document classifiers based on BoW and CBoW



recommended reading

- ▶ **Goldberg, chapter 8**: introduction to bag-of-words, embeddings, and continuous bag-of-words
- ▶ **Joulin et al. (2017)**: the fastText paper about text categorization

references I

- M. Iyyer, V. Manjunatha, J. Boyd-Graber, and H. Daumé III. 2015. [Deep unordered composition rivals synt. methods for text classification](#). In *ACL*.
- A. Joulin, E. Grave, P. Bojanowski, and T. Mikolov. 2017. [Bag of tricks for efficient text classification](#). In *EACL*.