

Deep Learning for Natural Language Processing

Training a word embedding model with the SGNS algorithm



UNIVERSITY OF
GOTHENBURG

CHALMERS

WASP | WALLENBERG AI
AUTONOMOUS SYSTEMS
AND SOFTWARE PROGRAM

Richard Johansson

`richard.johansson@gu.se`

training word embedding models on raw text

- ▶ we will now see how to train word embeddings from raw text
- ▶ according to the **distributional hypothesis**, word meaning is reflected in the distribution of contexts
- ▶ the key idea is to build a model of **cooccurrence**: for a word, what contexts are we likely to see?

training word embedding models on raw text

- ▶ we will now see how to train word embeddings from raw text
- ▶ according to the **distributional hypothesis**, word meaning is reflected in the distribution of contexts
- ▶ the key idea is to build a model of **cooccurrence**: for a word, what contexts are we likely to see?
- ▶ for instance: for *coffee*
 - ▶ *drink, black, cup* are likely contexts
 - ▶ *mosquito, purple, gradient* are unlikely contexts

skip-gram with negative sampling (word2vec)

- ▶ the **skip-gram with negative sampling** (SGNS) model is a well-known training method ([Mikolov et al., 2013](#))
- ▶ it is a simplification of several previous models
- ▶ in this model, we have one set of vectors V_T for the target words, and another set of vectors V_C for the contexts
- ▶ SGNS trains a model similar to logistic regression so that
 - ▶ $V_T(\text{coffee}) \cdot V_C(\text{drink})$ is high
 - ▶ $V_T(\text{coffee}) \cdot V_C(\text{gradient})$ is low

SGNS: the model

- ▶ collect **word-context pairs** occurring in the text data

| | | | | | | | | | |
|-----|-------|-------|------------------------------|--------------------------|--------------------|---|------------------------------|--|--|
| The | quick | brown | fox jumps over the lazy dog. | | | → | (the, quick) (the, brown) | | |
| The | quick | brown | fox | jumps over the lazy dog. | | | → | (quick, the) (quick, brown) (quick, fox) | |
| The | quick | brown | fox | jumps | over the lazy dog. | | | → | (brown, the) (brown, quick) (brown, fox) (brown, jumps) |

[[source](#)]

- ▶ for each pair, randomly generate a number of **synthetic pairs**:
fox: mechanic
fox: nitrogen
fox: persuade

SGNS: the model

- ▶ then fit the following model with respect to (V_T, V_C)

$$P(\text{true pair} | (w, c)) = \frac{1}{1 + \exp(-V_T(w) \cdot V_C(c))}$$


$$P(\text{synthetic pair} | (w, c)) = 1 - \frac{1}{1 + \exp(-V_T(w) \cdot V_C(c))}$$

- ▶ so the whole training objective becomes

$$\begin{aligned} \mathcal{L}(V_T, V_C) = & - \sum_{(w,c) \in P} \log \sigma(V_T(w) \cdot V_C(c)) \\ & - \sum_{(w,c) \in N} \log \sigma(-V_T(w) \cdot V_C(c)) \end{aligned}$$

a challenge when training word embedding models

- ▶ word distributions are highly skewed



- ▶ remember **Zipf's law**

engineering tricks in SGNS: removing highly frequent words

- ▶ a word occurrence is removed from the input with a frequency-dependent probability:

$$P_{\text{remove}}(w) = \max(0, 1 - \sqrt{\frac{t}{\text{freq}(w)}})$$


- ▶ reduces the influence of frequent words and speeds up training

engineering tricks in SGNS: negative sample distribution


- ▶ SGNS uses the following distribution for drawing negative contexts:

$$P_{\text{neg}}(c) \propto \text{freq}(c)^\alpha$$

- ▶ α is set to a number between 0 and 1




$\alpha = 1$



$\alpha = 0$

engineering tricks in SGNS: width of context window

- ▶ positive pairs are formed from words appearing in a window



- ▶ each time, a window width is drawn from a uniform distribution
- ▶ idea: distant words should be less influential

hyperparameters in SGNS

- ▶ **dimensionality** of embeddings: typically 50–300
- ▶ **number of negative samples**: typically 5
- ▶ **maximal width of the context window**: typically 5
- ▶ **smoothing constant** α for negative samples: typically 0.75
- ▶ **pruning threshold** t for frequent words: typically 0.0001

implementations

- ▶ **word2vec**: the software by Mikolov when he was at Google
 - ▶ implements the SGNS model (and a few others)
 - ▶ includes a model built by Google using a huge collection of news text
- ▶ **gensim**: a nice Python library by Řehůřek
 - ▶ includes a reimplementation of SGNS but also several other useful algorithms, such as LSA and LDA
 - ▶ the library also includes many pre-trained models

example notebooks

- ▶ using the Gensim library: training and loading models
- ▶ full implementation of the SGNS in PyTorch

references I

T. Mikolov, I. Sutskever, K. Chen, G. S. Corrado, and J. Dean. 2013.
[Distributed representations of words and phrases and their compositionality.](#)
In *NIPS*.