

Deep Learning for Natural Language Processing

More training methods for word embeddings



UNIVERSITY OF
GOTHENBURG

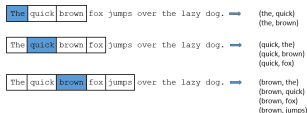
CHALMERS

WASP | WALLENBERG AI
AUTONOMOUS SYSTEMS
AND SOFTWARE PROGRAM

Richard Johansson

`richard.johansson@gu.se`

- ▶ while SGNS is probably the most well-known algorithm for training word embedding models, there are several others
- ▶ we'll take a quick tour of some of the most prominent
 - ▶ **“prediction-based”**: collecting training instances from individual occurrences (like SGNS)



- ▶ **“count-based”**: methods based on co-occurrence matrices

$$X = \begin{matrix} & \begin{matrix} I & like & enjoy & deep & learning & NLP & flying & . \end{matrix} \\ \begin{matrix} I \\ like \\ enjoy \\ deep \\ learning \\ NLP \\ flying \\ . \end{matrix} & \begin{bmatrix} 0 & 2 & 1 & 0 & 0 & 0 & 0 & 0 \\ 2 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 1 & 1 & 1 & 0 \end{bmatrix} \end{matrix}$$

continuous bag-of-words for training embeddings

- ▶ the continuous bag-of-words (CBoW) model considers the **whole context** instead of breaking it up into separate pairs:

*the quick brown **fox** jumps over the lazy dog*



*{ the, quick, brown, jumps, over, the }, **fox***

- ▶ the model is almost like SGNS:

$$P(\text{true pair} | (w, C)) = \frac{1}{1 + \exp(-V_T(w) \cdot V_C(C))}$$

where $V'(C)$ is the sum of context embeddings

$$V_C(C) = \sum_{c \in C} V_C(c)$$

- ▶ also available in the word2vec software

how can we deal with **out-of-vocabulary** words?

- ▶ what if *dingo* is in the vocabulary but not *dingoes*?
- ▶ humans can handle these kinds of situations!
- ▶ **fastText** (Bojanowski et al., 2017) modifies the SGNS model to handle these situations:

$$V_T(w) = \sum_{g \in \mathcal{G}} \mathbf{z}_g$$


where \mathcal{G} is the set of subwords for w :

$$\mathcal{G} = \{ \text{'<dingoes>'}, \text{'<di'}, \text{'din'}, \text{'ing'}, \dots, \text{'ngoes>'} \}$$

- ▶ handles rare words and OOV words better than SGNS

how can we deal with **multiple senses**?


- ▶ for instance, **rock** can refer to a stone or a type of music



- ▶ Luis's PhD thesis from a few weeks back: *Splitting rocks: Learning word sense representations from corpora and lexica*

compare: matrix factorization in recommender systems


- the most famous approach in recommenders is based on **factorization** of the user/item rating matrix



- to predict a missing cell (rating of an unseen item):

$$\hat{r}_{ui} = \mathbf{p}_u \cdot \mathbf{q}_i$$

where \mathbf{p}_u is the user's vector, and \mathbf{q}_i the item's vector



example of a word–word co-occurrence matrix

- ▶ assume we have the following set of texts:

- ▶ “I like NLP”
- ▶ “I like deep learning”
- ▶ “I enjoy flying”

$$X = \begin{array}{c} \begin{matrix} & I & like & enjoy & deep & learning & NLP & flying & . \end{matrix} \\ \begin{matrix} I \\ like \\ enjoy \\ deep \\ learning \\ NLP \\ flying \\ . \end{matrix} \end{array} \begin{bmatrix} 0 & 2 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2 & 0 & 0 & 1 & 0 & 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 1 & 1 & 1 & 1 & 0 \end{bmatrix}$$

[[source](#)]

matrix-based word embeddings


- ▶ **Latent Semantic Analysis** (Landauer and Dumais, 1997) was the first vector-based word representation model
 - ▶ it applies singular value decomposition (SVD) to a word–document matrix
- ▶ several other approaches:
 - ▶ different types of matrices (word–document, word–word, ...)
 - ▶ transformations of the matrix (log, PMI, ...)
 - ▶ different types of factorizations (none, SVD, NMF, ...)
- ▶ Levy and Goldberg (2014) show a connection between SGNS and matrix-based methods

GloVe

- ▶ GloVe (Pennington et al., 2014) is a famous matrix-based method
 - ▶ <https://nlp.stanford.edu/projects/glove/>
- ▶ they try to find word and context vectors to satisfy

$$V_T(w) \cdot V_C(c) + b_T(w) + b_C(c) = \log \#(w, c)$$

- ▶ they downweight infrequent words



what should we prefer, count-based or prediction-based?

- ▶ see [Baroni et al. \(2014\)](#) for a comparison of count-based and prediction-based
 - ▶ they come out strongly in favor of prediction-based
 - ▶ but this result has been questioned
- ▶ pros and cons:
 - ▶ prediction-based methods are sensitive to the order the examples are processed
 - ▶ count-based methods can be messy to implement with a large vocabulary

references I

- M. Baroni, G. Dinu, and G. Kruszewski. 2014. Don't count, predict! A systematic comparison of context-counting vs. context-predicting semantic vectors. In *ACL*.
- P. Bojanowski, E. Grave, A. Joulin, and T. Mikolov. 2017. Enriching word vectors with subword information. *TACL* 5:135–146.
- T. K. Landauer and S. T. Dumais. 1997. A solution to Plato's problem: The latent semantic analysis theory of acquisition, induction and representation of knowledge. *Psychological Review* 104:211–240.
- O. Levy and Y. Goldberg. 2014. Neural word embedding as implicit matrix factorization. In *NIPS*.
- L. Nieto Piña. 2019. *Splitting rocks: Learning word sense representations from corpora and lexica*. Ph.D. thesis, University of Gothenburg.
- J. Pennington, R. Socher, and C. Manning. 2014. GloVe: Global vectors for word representation. In *EMNLP*.