

# Deep Learning for Natural Language Processing

## The Transformer model



UNIVERSITY OF  
GOTHENBURG

---


**CHALMERS**

**WASP** | WALLENBERG AI  
AUTONOMOUS SYSTEMS  
AND SOFTWARE PROGRAM

**Richard Johansson**

`richard.johansson@gu.se`


# drawbacks of recurrent models



- ▶ even with GRUs and LSTMs, it is difficult for RNNs to preserve information over long distances
- ▶ we introduced **attention** as a way to deal with this problem
- ▶ can we skip the RNN and **just use attention**?

## attention models: recap

- first, compute an “energy”  $e_i$  for each state  $\mathbf{h}_i$



- for the attention weights, we apply the softmax:


$$\alpha_i = \frac{\exp e_i}{\sum_{j=1}^n \exp e_j}$$

- finally, the “summary” is computed as a weighted sum

$$\mathbf{s} = \sum_{i=1}^n \alpha_i \mathbf{h}_i$$

# the Transformer

- ▶ the Transformer (Vaswani et al., 2017) is an architecture that uses attention for information flow  
*“Attention is all you need”*
- ▶ it was originally designed for machine translation and has two parts:
  - ▶ an **encoder** that “summarizes” an input sentence
  - ▶ a **decoder** (a conditional LM) that generates an output, based on the input
- ▶ let's consider the encoder




# illustration of a Transformer block

**Jane lives in Paris . She is a dentist .**


# illustration of a Transformer block

Jane lives in Paris . She is a dentist .

# illustration of a Transformer block




# illustration of a Transformer block

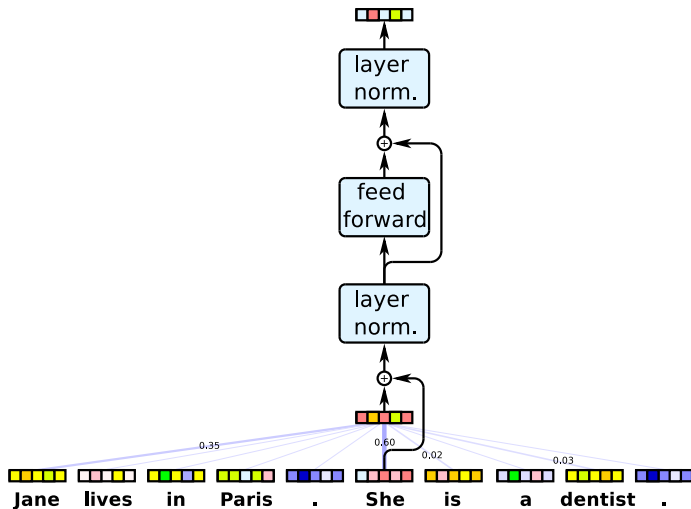




# illustration of a Transformer block



# illustration of a Transformer block




# multi-head attention

- ▶ in each layer, the Transformer applies several attention models (“heads”) in parallel
- ▶ intuitively, the heads are “looking” for different types of information
- ▶ each attention head computes a **scaled dot product attention**:


$$e_i = \frac{1}{\sqrt{d}} \mathbf{q}_i \cdot \mathbf{k}_j$$
$$\alpha = \text{softmax}(\mathbf{e})$$

where  $\mathbf{q}_i$  and  $\mathbf{k}_j$  are linear transformations of the input at positions  $i$  and  $j$

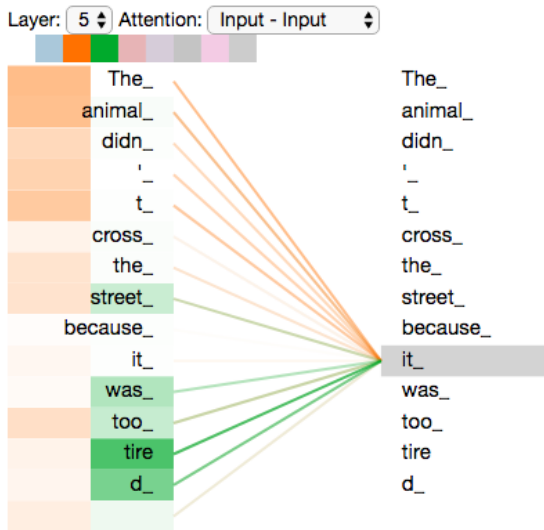


## a layer in the Transformer encoder

- ▶ after each application of multi-head attention, a 2-layer feedforward model (with ReLU activation) is applied
- ▶ residual connections (“shortcuts”) and layer normalization (Ba et al., 2016) added for robustness and to facilitate training
- ▶ the Transformer encoder consists of a stack of this type of block



# what do the attention heads look at?



► see (Vig, 2019)

# pros and cons

- + short path length for information flow
- quadratic complexity

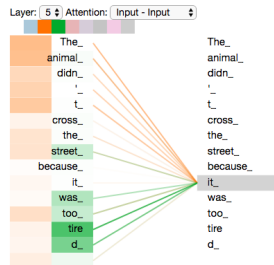
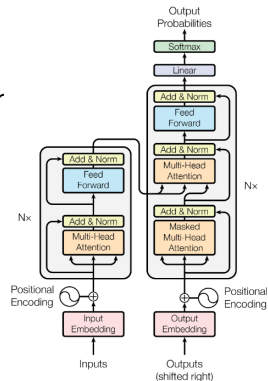


Table 1: Maximum path lengths, per-layer complexity and minimum number of sequential operations for different layer types.  $n$  is the sequence length,  $d$  is the representation dimension,  $k$  is the kernel size of convolutions and  $r$  the size of the neighborhood in restricted self-attention.

Layer Type	Complexity per Layer	Sequential Operations	Maximum Path Length
Self-Attention	$O(n^2 \cdot d)$	$O(1)$	$O(1)$
Recurrent	$O(n \cdot d^2)$	$O(n)$	$O(n)$
Convolutional	$O(k \cdot n \cdot d^2)$	$O(1)$	$O(\log_k(n))$
Self-Attention (restricted)	$O(r \cdot n \cdot d)$	$O(1)$	$O(n/r)$

# the road ahead

- ▶ the full Transformer is an effective model for machine translation
- ▶ we'll return to it when we discuss **encoder-decoder** architectures
- ▶ for now, let's use it simply as a pre-trained representation



## The Illustrated Transformer



# references

- J. Ba, J. Kiros, and G. Hinton. 2016. [Layer normalization](#). arXiv:1607.06450.
- A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin. 2017. [Attention is all you need](#). In *NIPS 30*.
- J. Vig. 2019. [Visualizing attention in transformer-based language representation models](#). arXiv:1904.02679.