

# Deep Learning for Natural Language Processing

## Factorized Sequence Models



UNIVERSITY OF  
GOTHENBURG

---

**CHALMERS**

**WASP** | WALLENBERG AI  
AUTONOMOUS SYSTEMS  
AND SOFTWARE PROGRAM

**Richard Johansson**

`richard.johansson@gu.se`

# Algorithmic approaches

- **Exhaustive search**

Cast structured prediction as a combinatorial optimisation problem over the set of target representations.

Viterbi algorithm, Eisner algorithm

- **Greedy search**

Cast structured prediction as a sequence of classification problems: at each point in time, predict one of several options.

window-based part-of-speech tagging, arc-standard algorithm

# Algorithmic approaches

- **Exhaustive search**

Cast structured prediction as a combinatorial optimisation problem over the set of target representations.

Viterbi algorithm, Eisner algorithm

- **Greedy search**

Cast structured prediction as a sequence of classification problems: at each point in time, predict one of several options.

window-based part-of-speech tagging, arc-standard algorithm

## general approach

- ▶ we define a scoring function over the **whole sequence**
- ▶ we maximize this function **over all possible sequences**

$$\hat{y} = \underset{y}{\operatorname{argmax}} \operatorname{score}(x, y; \theta)$$

Diagram illustrating the components of the scoring function:

- $\hat{y}$  is labeled "predicted output".
- $x$  is labeled "input".
- $y$  is labeled "candidate output".
- $\theta$  is labeled "model parameters".

Eisenstein (2019), § 1.2.2

there are many possible sequences. . .

I	want	to	live	in	peace
PRON	VERB	PART	VERB	ADP	NOUN
NOUN	NOUN	ADP	ADJ	ADV	VERB
		ADV	ADV	ADJ	
				NOUN	

# making the arg max problem tractable

- ▶ the number of possible sequences is typically **exponential** in the length of the input
- ▶ we will need to make **assumptions** about the scoring function
- ▶ then we can design special algorithms to compute the arg max

I	want	to	live	in	peace
PRON	VERB	PART	VERB	ADP	NOUN
NOUN	NOUN	ADP	ADJ	ADV	VERB
		ADV	ADV	ADJ	
				NOUN	

# first-order factorized scoring function

- ▶ we will work with scoring functions where we compute a sum over “parts” or “factors”:

$$\text{score}(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^L \phi^e(\mathbf{x}, y_i) + \sum_{i=1}^L \phi^t(\mathbf{x}, y_{i-1}, y_i)$$

- ▶ the part scoring functions  $\phi^e$  and  $\phi^t$  compute scores for single labels and pairs of adjacent labels, respectively
- ▶ it's a **first-order** factorization: we model 1-step interactions

## what are these scores?


$$\phi^e(\mathbf{x}, y_i) \quad \phi^t(\mathbf{x}, y_{i-1}, y_i)$$

- ▶ following HMM terminology, we call them “**emission** scores” ( $\phi^e$ ) and “**transition** scores” ( $\phi^t$ )
  - ▶ in a HMM, they are log probabilities
  - ▶ we will use neural networks to compute them instead
- ▶ in the next lecture, we'll discuss how to train them
- ▶ for now, let's focus on **decoding**: how to maximize  $\text{score}(\mathbf{x}, \mathbf{y})$




# first-order factorization as a graph

- ▶ we can construct a **graph** where each node represents a possible label
- ▶ we compute node scores with  $\phi^e$  and edge scores with  $\phi^t$



# first-order factorization as a graph


- ▶ we can construct a **graph** where each node represents a possible label
- ▶ we compute node scores with  $\phi^e$  and edge scores with  $\phi^t$



- ▶ maximizing the scoring function is equivalent to finding the **highest-scoring path**

# the Viterbi algorithm


- ▶ the **Viterbi algorithm** can be used to find the highest-scoring path in this type of graph
- ▶ you may recognize this as a special case of the **max-sum algorithm** for graphical models



- ▶ it is a **dynamic programming** algorithm
  - ▶ it proceeds from left to right
  - ▶ the optimal paths in step  $i$  are computed by considering the optimal paths in step  $i - 1$


# dynamic programming in the Viterbi algorithm

- ▶ to compute the best path ending with *saw* as a verb, **consider the best paths for the previous word** and the **transition scores**  $\phi^t(\mathbf{x}, y_{i-1}, y_i)$





# dynamic programming in the Viterbi algorithm

- ▶ to compute the best path ending with *saw* as a verb, **consider the best paths for the previous word** and the **transition scores**  $\phi^t(\mathbf{x}, y_{i-1}, y_i)$





# dynamic programming in the Viterbi algorithm

- to compute the best path ending with *saw* as a verb, **consider the best paths for the previous word** and the **transition scores**  $\phi^t(\mathbf{x}, y_{i-1}, y_i)$



# dynamic programming in the Viterbi algorithm

- to compute the best path ending with *saw* as a verb, **consider the best paths for the previous word** and the **transition scores**  $\phi^t(\mathbf{x}, y_{i-1}, y_i)$




# Viterbi example


<S>      I      can      see      a      can      .      <E>




# Viterbi example




# Viterbi example




# Viterbi example




# Viterbi example




# Viterbi example




# Viterbi example



# Viterbi example




# Viterbi example







# Viterbi example



# Viterbi example



# Viterbi example



**next up:** how to train the scoring functions  $\phi^e$  and  $\phi^t$